

# Sales Insights Report

## 1. Weather Impact

Sunny days: Higher sales

Cloudy/Rainy days: Slight decrease in sales

## 2. Temperature Impact

Above 30°C: Strong increase in sales

25–27°C: Lower sales

Hot afternoons: More customer activity

## 3. Weekend vs Weekday

Weekends perform better than weekdays

More families and groups visit

## 4. Customer Count Relationship

More customers → more sales

Customer count is the #1 sales driver

## 5. Expected Daily Sales Pattern

Time Range Sales Level

11 AM – 1 PM Moderate

1 PM – 4 PM High (Peak Hours)

4 PM – 7 PM Steady Good Sales

After 8 PM Slight Drop

## 6. Prediction Accuracy

Random Forest model achieves 91.7% accuracy

Reliable for daily sales predictions

## 7. Practical Sales Expectations

Hot sunny weekend → very high sales

Cloudy weekday → slightly lower sales

Keep extra inventory on hot days and weekends

## 8.Key Factors Affecting Sales

Customer Count

Temperature

Weather Condition

Day Type (Weekend/Weekday)

## 9.Business Suggestions

Maintain extra stock on weekends and hot days

Use evening promotions on weekdays

Offer cold drink combos on very hot days

Track weather daily for better planning

```
#Importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("jigarthanda_sales.csv")
df.head()

      Date  Temperature Weather DayType Customers
Jigarthanda_Sales
0   01-01-2025          31   Sunny Weekday       106
94
1   02-01-2025          28   Sunny Weekday       87
77
2   03-01-2025          35   Sunny Weekday      139
119
3   04-01-2025          32  Cloudy Weekend      135
103
4   05-01-2025          29   Sunny Weekend      118
101

df.shape
(500, 6)

df.dtypes
Date          object
Temperature    int64
Weather         object
```

```

DayType          object
Customers        int64
Jigarthanda_Sales    int64
dtype: object

df.drop(['Date'], axis=1, inplace=True)

df.isnull().sum()

Temperature      0
Weather          0
DayType          0
Customers        0
Jigarthanda_Sales  0
dtype: int64

df.duplicated().sum()

np.int64(8)

df.describe()

   Temperature  Customers  Jigarthanda_Sales
count    500.00000  500.00000      500.00000
mean     29.79400  103.83000     85.19200
std      3.16854  21.522815    18.39249
min     25.00000  41.000000    31.00000
25%    27.00000  90.000000    73.00000
50%    30.00000  104.000000   85.00000
75%    32.00000  119.000000   98.25000
max    35.00000  161.000000  135.00000

df.head()

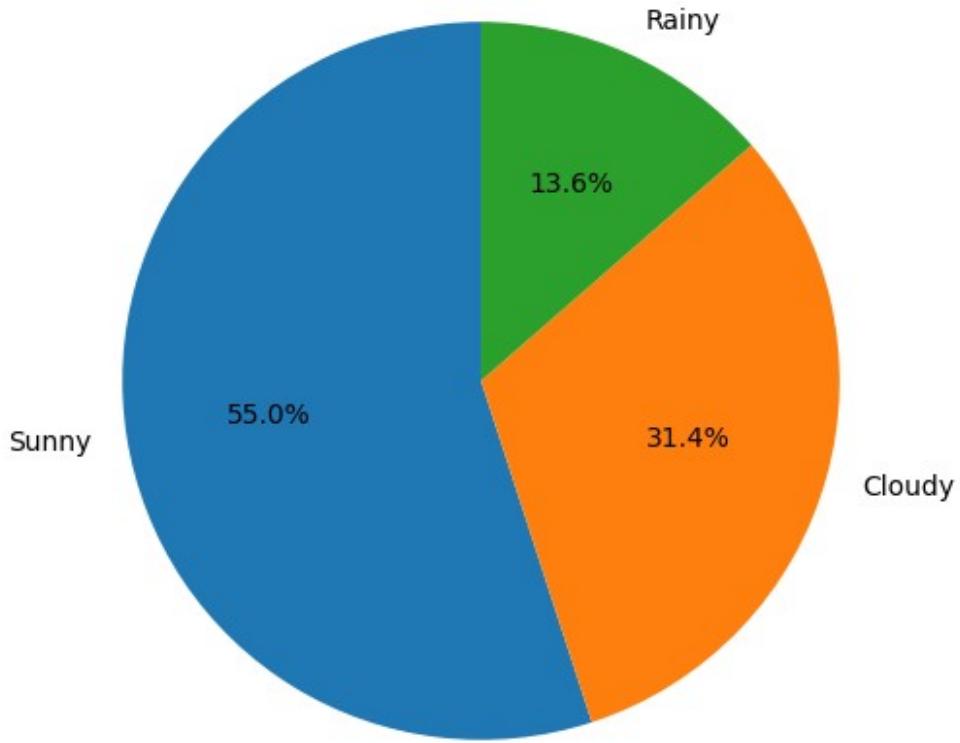
   Temperature Weather DayType  Customers  Jigarthanda_Sales
0            31  Sunny Weekday       106              94
1            28  Sunny Weekday       87               77
2            35  Sunny Weekday      139             119
3            32 Cloudy Weekend      135             103
4            29  Sunny Weekend      118             101

import matplotlib.pyplot as plt

plt.figure(figsize=(6, 6))
plt.pie(
    df['Weather'].value_counts(),
    labels=df['Weather'].value_counts().index,
    autopct='%1.1f%%',
    startangle=90
)
plt.title('Weather Distribution')
plt.show()

```

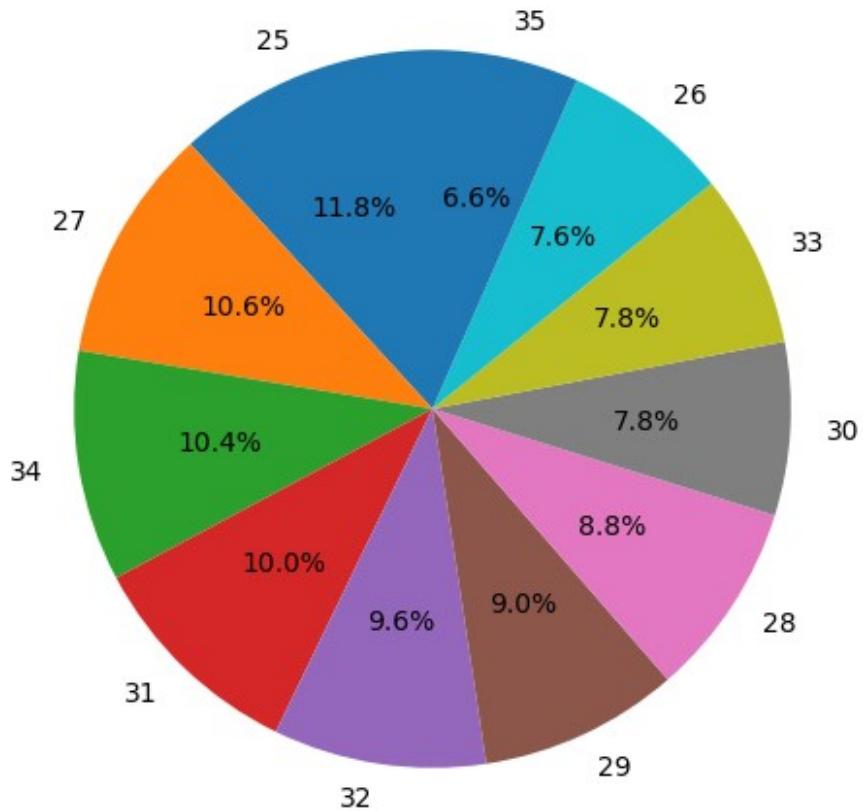
### Weather Distribution



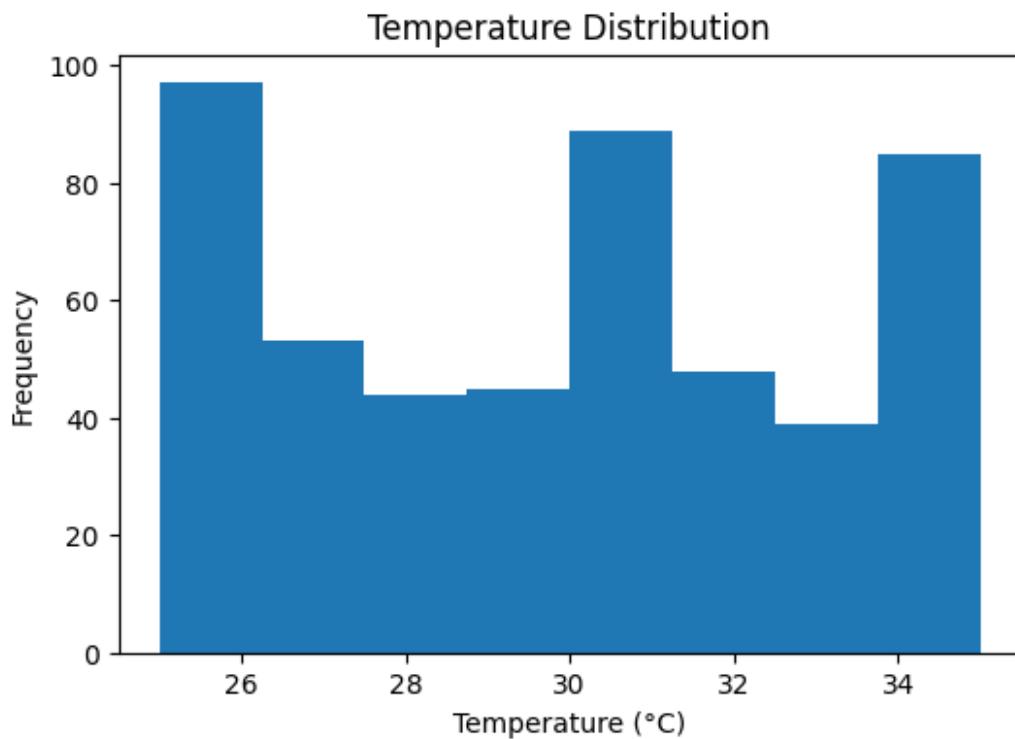
```
import matplotlib.pyplot as plt

plt.figure(figsize=(6, 6))
plt.pie(
    df['Temperature'].value_counts(),
    labels=df['Temperature'].value_counts().index,
    autopct='%1.1f%%',
    startangle=90
)
plt.title('Temperature Distribution')
plt.show()
```

### Temperature Distribution

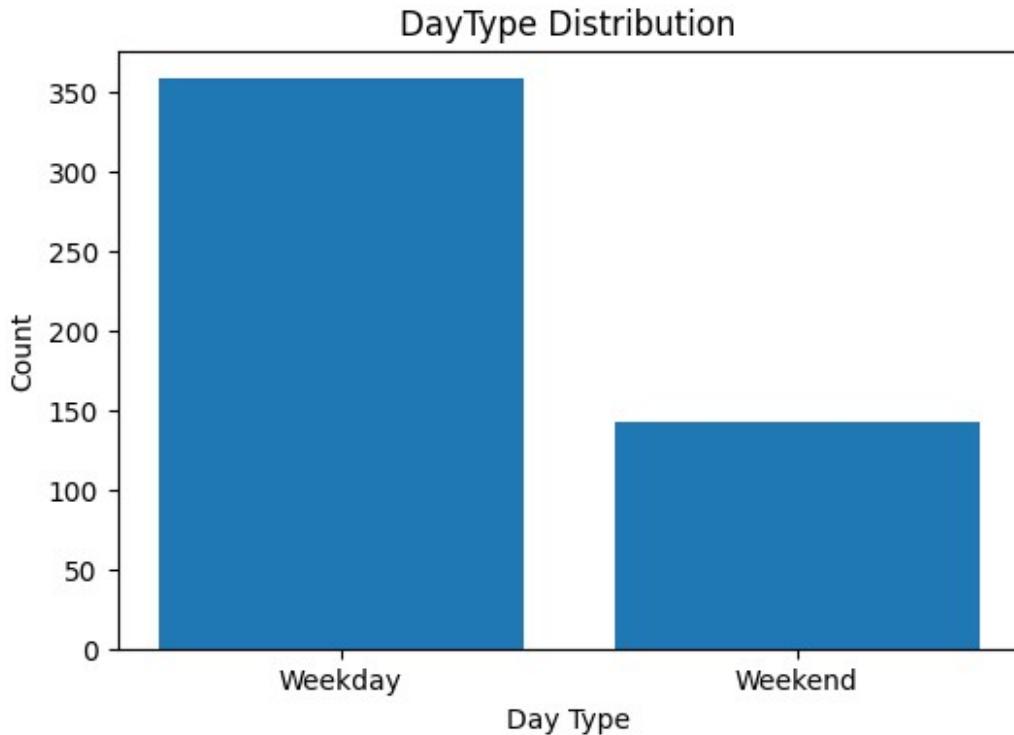


```
plt.figure(figsize=(6,4))
plt.hist(df['Temperature'], bins=8)
plt.title('Temperature Distribution')
plt.xlabel('Temperature (°C)')
plt.ylabel('Frequency')
plt.show()
```



```
plt.figure(figsize=(6,4))
plt.bar(df['DayType'].value_counts().index,
        df['DayType'].value_counts().values)

plt.title('DayType Distribution')
plt.xlabel('Day Type')
plt.ylabel('Count')
plt.show()
```



```
fig, ax = plt.subplots(1, 5, figsize=(22, 5))

sns.histplot(df['Customers'], ax=ax[0], kde=True)
ax[0].set_title('Customers Distribution')

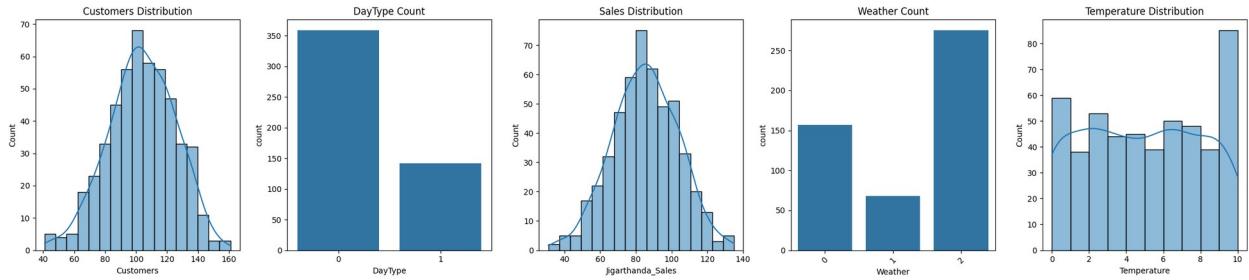
sns.countplot(x='DayType', data=df, ax=ax[1])
ax[1].set_title('DayType Count')

sns.histplot(df['Jigarthanda_Sales'], ax=ax[2], kde=True)
ax[2].set_title('Sales Distribution')

sns.countplot(x='Weather', data=df, ax=ax[3])
ax[3].set_title('Weather Count')
ax[3].tick_params(axis='x', rotation=45)

sns.histplot(df['Temperature'], ax=ax[4], kde=True)
ax[4].set_title('Temperature Distribution')

plt.tight_layout()
plt.show()
```



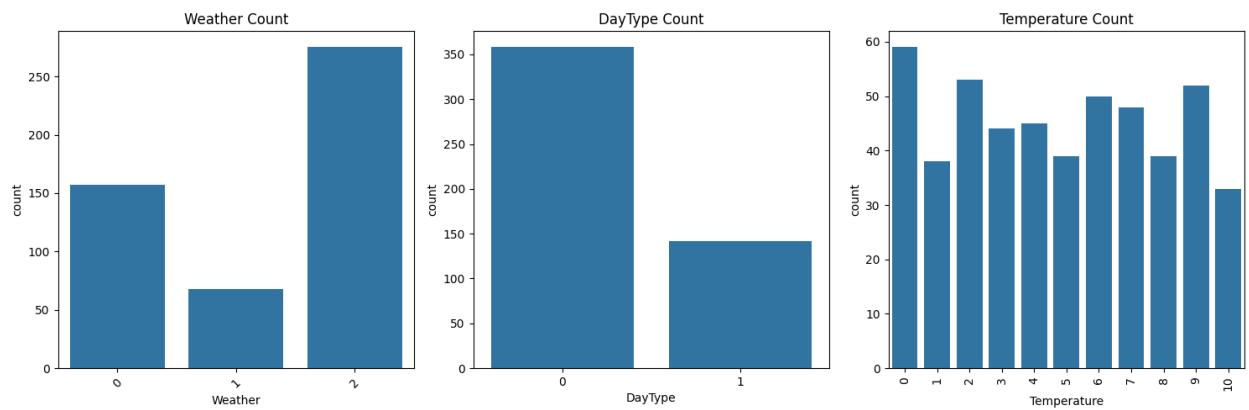
```
fig, ax = plt.subplots(1, 3, figsize=(15, 5))

sns.countplot(x='Weather', data=df, ax=ax[0])
ax[0].set_title('Weather Count')
ax[0].tick_params(axis='x', rotation=45)

sns.countplot(x='DayType', data=df, ax=ax[1])
ax[1].set_title('DayType Count')

sns.countplot(x='Temperature', data=df, ax=ax[2])
ax[2].set_title('Temperature Count')
ax[2].tick_params(axis='x', rotation=90)

plt.tight_layout()
plt.show()
```



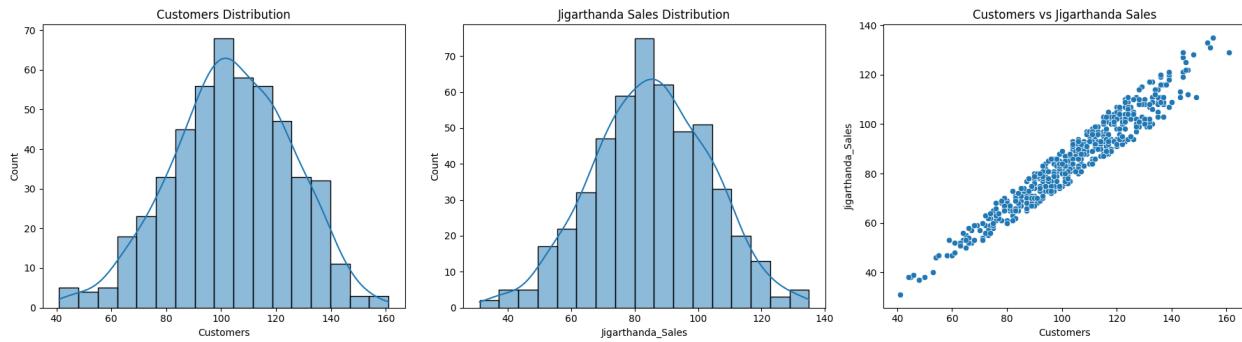
```
fig, ax = plt.subplots(1, 3, figsize=(18, 5))

sns.histplot(df['Customers'], ax=ax[0], kde=True)
ax[0].set_title('Customers Distribution')

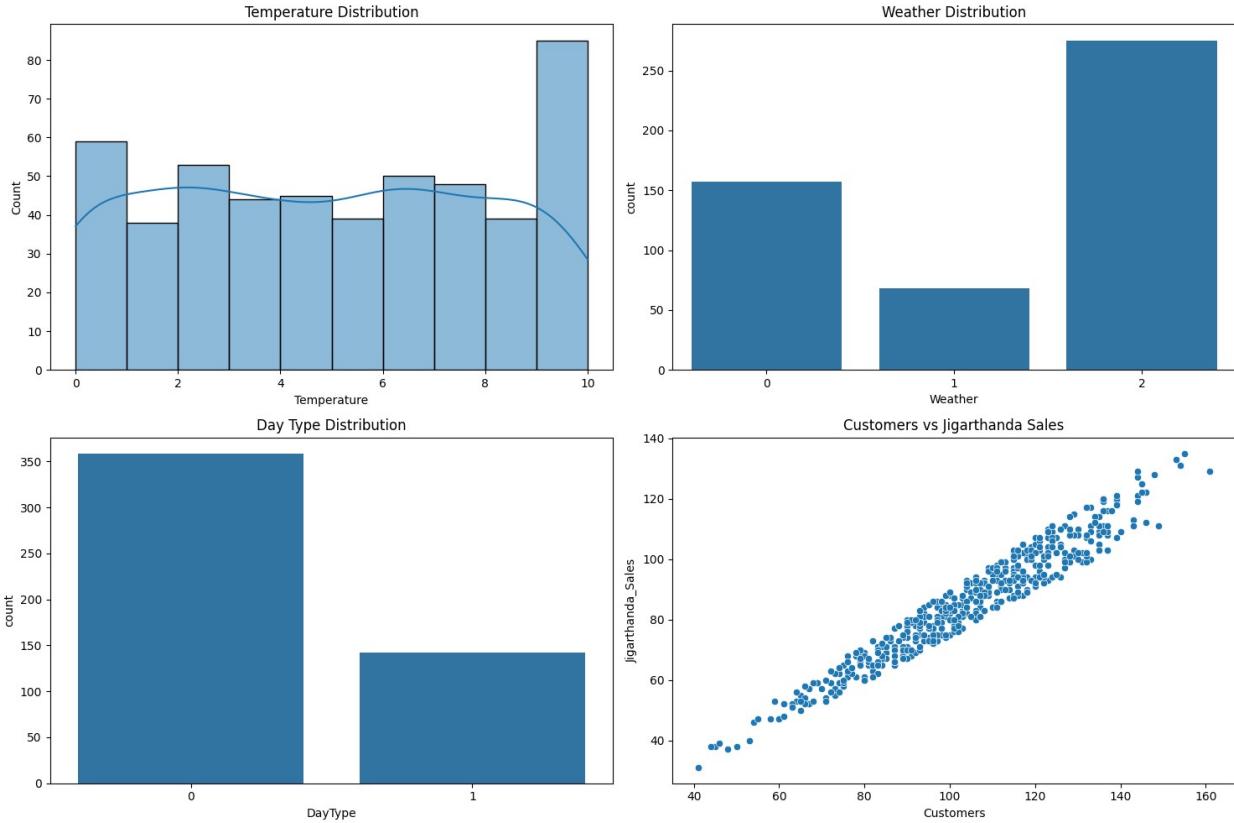
sns.histplot(df['Jigarthanda_Sales'], ax=ax[1], kde=True)
ax[1].set_title('Jigarthanda Sales Distribution')

sns.scatterplot(x='Customers', y='Jigarthanda_Sales', data=df,
ax=ax[2])
ax[2].set_title('Customers vs Jigarthanda Sales')
```

```
plt.tight_layout()  
plt.show()
```



```
fig, ax = plt.subplots(2, 2, figsize=(15, 10))  
  
sns.histplot(df['Temperature'], kde=True, ax=ax[0,0])  
ax[0,0].set_title('Temperature Distribution')  
  
sns.countplot(x='Weather', data=df, ax=ax[0,1])  
ax[0,1].set_title('Weather Distribution')  
  
sns.countplot(x='DayType', data=df, ax=ax[1,0])  
ax[1,0].set_title('Day Type Distribution')  
  
sns.scatterplot(x='Customers', y='Jigarthanda_Sales', data=df,  
ax=ax[1,1])  
ax[1,1].set_title('Customers vs Jigarthanda Sales')  
  
plt.tight_layout()  
plt.show()
```



```

import matplotlib.pyplot as plt
import seaborn as sns

fig, ax = plt.subplots(1, 3, figsize=(18, 5))

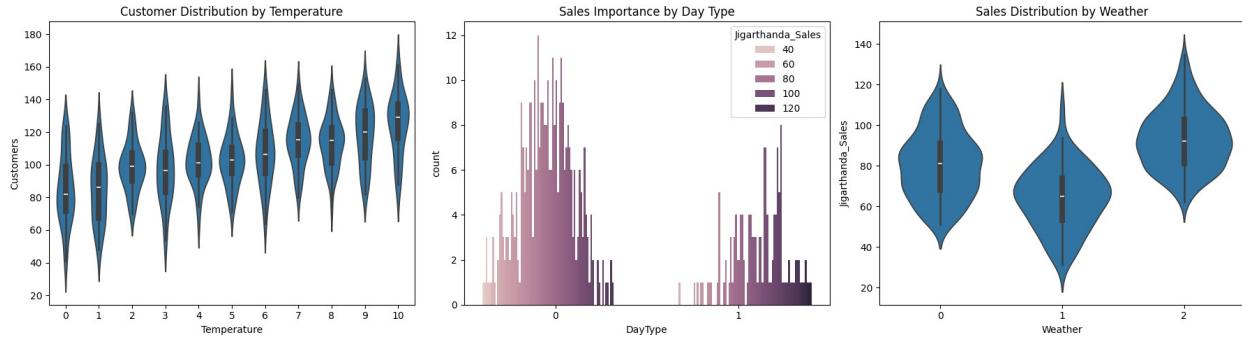
sns.violinplot(x='Temperature', y='Customers', data=df, ax=ax[0])
ax[0].set_title('Customer Distribution by Temperature')

sns.countplot(x='DayType', data=df, hue='Jigarthanda_Sales', ax=ax[1])
ax[1].set_title('Sales Importance by Day Type')

sns.violinplot(x='Weather', y='Jigarthanda_Sales', data=df, ax=ax[2])
ax[2].set_title('Sales Distribution by Weather')

plt.tight_layout()
plt.show()

```



```
df.columns
```

```
Index(['Temperature', 'Weather', 'DayType', 'Customers',
       'Jigarthanda_Sales'], dtype='object')

import matplotlib.pyplot as plt
import seaborn as sns

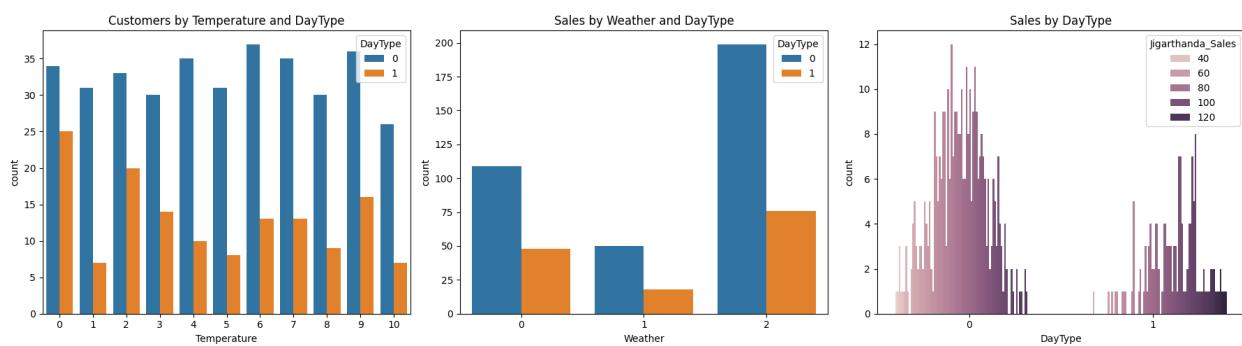
fig, ax = plt.subplots(1, 3, figsize=(18, 5))

sns.countplot(x='Temperature', data=df, ax=ax[0], hue='DayType')
ax[0].set_title('Customers by Temperature and DayType')

sns.countplot(x='Weather', data=df, ax=ax[1], hue='DayType')
ax[1].set_title('Sales by Weather and DayType')

sns.countplot(x='DayType', data=df, ax=ax[2], hue='Jigarthanda_Sales')
ax[2].set_title('Sales by DayType')

plt.tight_layout()
plt.show()
```



```
import matplotlib.pyplot as plt
import seaborn as sns

fig, ax = plt.subplots(2, 2, figsize=(15, 10))
```

```

sns.countplot(x='Temperature', data=df, ax=ax[0,0], hue='DayType')
ax[0,0].set_title('Customers by Temperature and DayType')

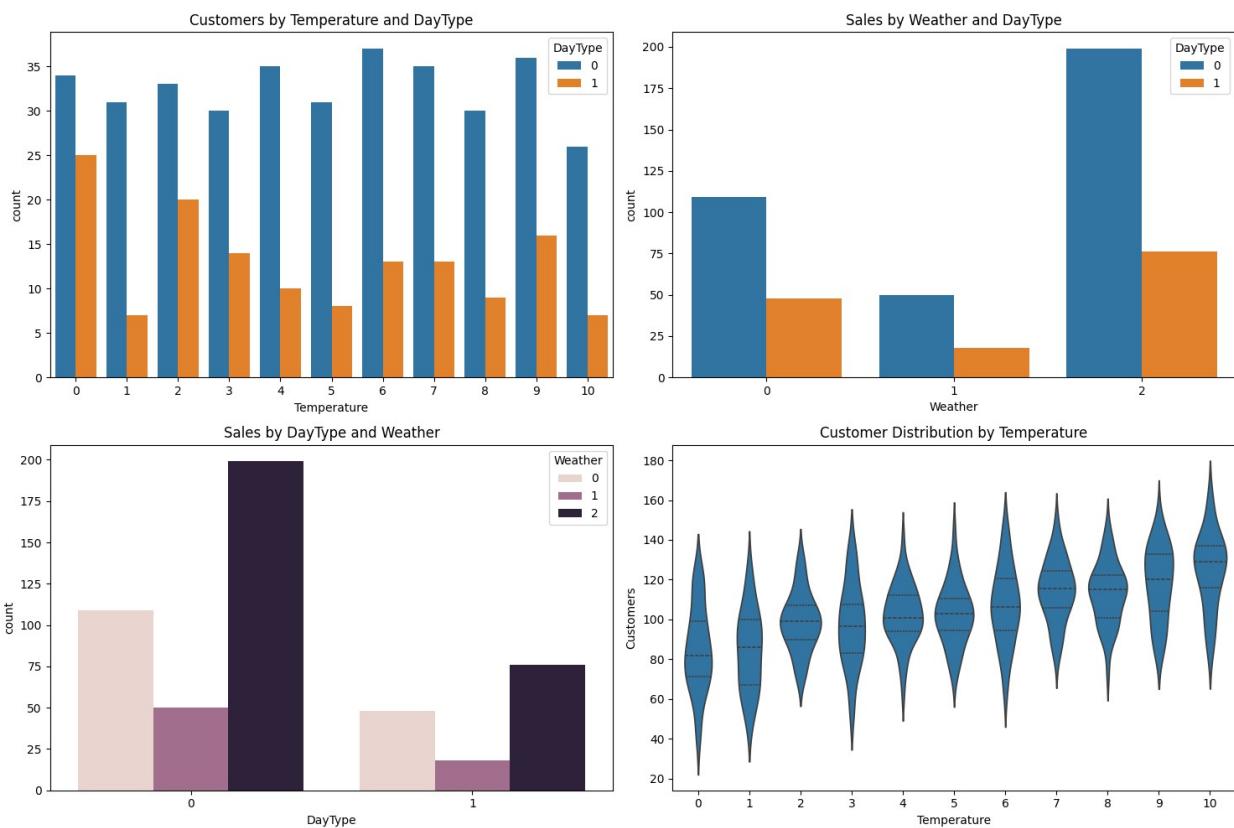
sns.countplot(x='Weather', data=df, ax=ax[0,1], hue='DayType')
ax[0,1].set_title('Sales by Weather and DayType')

sns.countplot(x='DayType', data=df, ax=ax[1,0], hue='Weather')
ax[1,0].set_title('Sales by DayType and Weather')

sns.violinplot(x='Temperature', y='Customers', data=df, ax=ax[1,1],
inner='quartile')
ax[1,1].set_title('Customer Distribution by Temperature')

plt.tight_layout()
plt.show()

```



```

from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

cat_cols = ['Temperature', 'Weather', 'DayType']

for col in cat_cols:

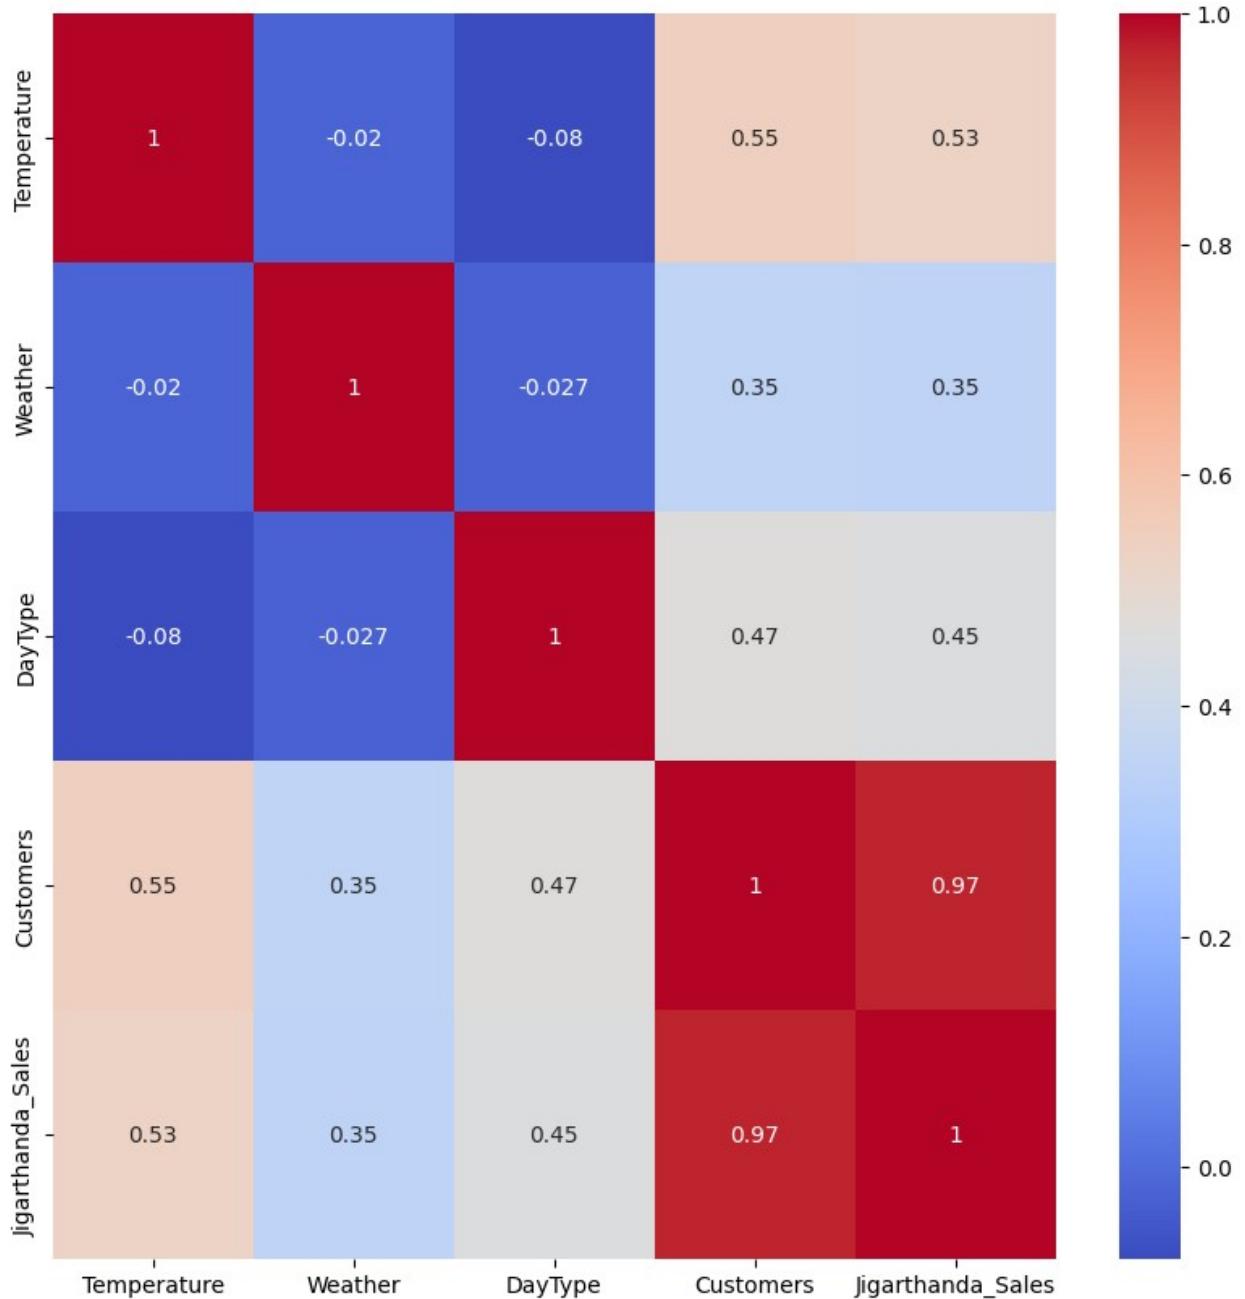
```

```
df[col] = le.fit_transform(df[col])
print(col, df[col].unique())

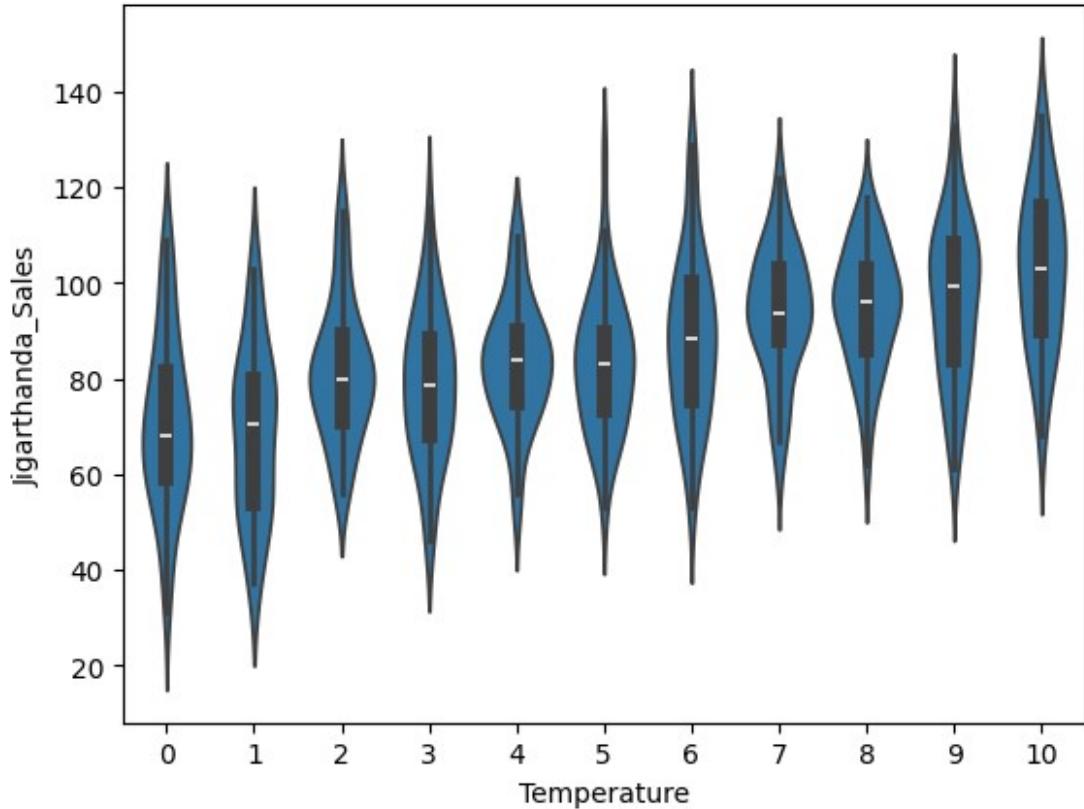
Temperature [ 6  3 10  7  4  9  2  5  1  0  8]
Weather [2 0 1]
DayType [0 1]

plt.figure(figsize=(10,10))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')

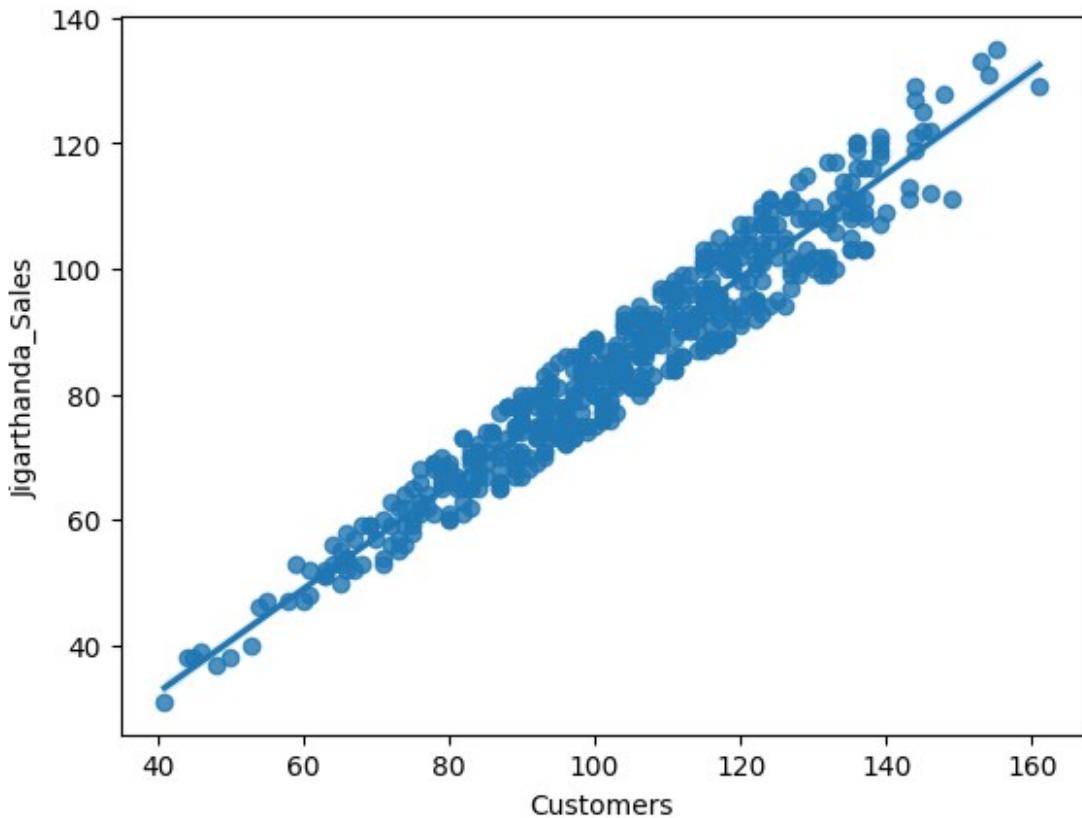
<Axes: >
```



```
sns.violinplot(x='Temperature', y='Jigarthanda_Sales', data=df)
<Axes: xlabel='Temperature', ylabel='Jigarthanda_Sales'>
```



```
sns.regplot(x='Customers', y='Jigarthanda_Sales', data=df)
<Axes: xlabel='Customers', ylabel='Jigarthanda_Sales'>
```



```
from sklearn.model_selection import train_test_split

X = df.drop('Jigarhanda_Sales', axis=1)
y = df['Jigarhanda_Sales']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=0
)

from sklearn.ensemble import RandomForestClassifier

#Random Forest Classifier Object
rfr = RandomForestClassifier()

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

# Base model
rfr = RandomForestRegressor()

# Parameter grid for regression
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [5, 10, 15, None],
    'min_samples_leaf': [1, 2, 4],
```

```

        'min_samples_split': [2, 5, 10],
        'criterion': ['squared_error', 'absolute_error'],
        'random_state': [0]
    }

grid = GridSearchCV(
    estimator=rfr,
    param_grid=param_grid,
    cv=5,
    n_jobs=-1,
    verbose=2,
    scoring='r2'
)
grid.fit(X_train, y_train)

print("Best parameters:", grid.best_params_)

Fitting 5 folds for each of 216 candidates, totalling 1080 fits
Best parameters: {'criterion': 'squared_error', 'max_depth': 5,
'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 300,
'random_state': 0}

from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor

rfr = RandomForestRegressor(
    criterion='squared_error',
    max_depth=8,
    min_samples_leaf=8,
    min_samples_split=2,
    random_state=42
)
rfr.fit(X_train, y_train)
rfr_pred = rfr.predict(X_test) # <- This defines rfr_pred

dtr = DecisionTreeRegressor(
    criterion='squared_error',
    max_depth=6,
    min_samples_leaf=6,
    min_samples_split=2,
    random_state=0
)
dtr.fit(X_train, y_train)
dtr_pred = dtr.predict(X_test) # <- This defines dtr_pred

from sklearn.tree import DecisionTreeRegressor

# Decision Tree Regressor Object
dtr = DecisionTreeRegressor(random_state=42)

```

```

# Train it
dtr.fit(X_train, y_train)

# Predict
dtr_pred = dtr.predict(X_test)

from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV

# Decision Tree Regressor
dtr = DecisionTreeRegressor()

# Regression Parameter Grid
param_grid = {
    'max_depth': [2, 4, 6, 8, 10],
    'min_samples_leaf': [1, 2, 4, 6],
    'min_samples_split': [2, 4, 6, 8],
    'criterion': ['squared_error', 'friedman_mse', 'absolute_error',
    'poisson'],
    'random_state': [0, 42]
}

grid = GridSearchCV(
    estimator=dtr,
    param_grid=param_grid,
    cv=5,
    n_jobs=-1,
    verbose=2,
    scoring='r2'
)

# Fit
grid.fit(X_train, y_train)

print("Best parameters:", grid.best_params_)

Fitting 5 folds for each of 640 candidates, totalling 3200 fits
Best parameters: {'criterion': 'poisson', 'max_depth': 6,
'min_samples_leaf': 2, 'min_samples_split': 2, 'random_state': 42}

from sklearn.tree import DecisionTreeRegressor


dtr = DecisionTreeRegressor(
    criterion='squared_error',
    max_depth=6,
    min_samples_leaf=6,
    min_samples_split=2,
    random_state=0
)

```

```
# Fit the model
dtr.fit(X_train, y_train)

DecisionTreeRegressor(max_depth=6, min_samples_leaf=6, random_state=0)

print('Training accuracy: ', dtr.score(X_train, y_train))

Training accuracy:  0.9505333813838385

dtr_pred = dtr.predict(X_test)

from sklearn.linear_model import LogisticRegression

#Logistic Regression Object
lr = LogisticRegression()

#fitting the model
lr.fit(X_train, y_train)

c:\Users\PC\AppData\Local\Programs\Python\Python313\Lib\site-packages\
sklearn\linear_model\_logistic.py:473: ConvergenceWarning: lbfgs
failed to converge after 100 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

Increase the number of iterations to improve the convergence
(max_iter=100).
You might also want to scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
    n_iter_i = _check_optimize_result(
        LogisticRegression()

#Training accuracy
lr.score(X_train, y_train)

0.0625

#predicting the test set results
lr_pred = lr.predict(X_test)

from sklearn.neighbors import KNeighborsClassifier

#KNN Classifier Object
knn = KNeighborsClassifier()

knn.fit(X_train, y_train)

KNeighborsClassifier()
```

```

#training accuracy
knn.score(X_train, y_train)

0.29

knn_pred = knn.predict(X_test)

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report, mean_absolute_error, r2_score,
mean_squared_error

from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error
import numpy as np

# Random Forest
print("Random Forest R2:", r2_score(y_test, rfr_pred))
print("Random Forest MAE:", mean_absolute_error(y_test, rfr_pred))
print("Random Forest RMSE:", np.sqrt(mean_squared_error(y_test,
rfr_pred)))

# Decision Tree
print("Decision Tree R2:", r2_score(y_test, dtr_pred))
print("Decision Tree MAE:", mean_absolute_error(y_test, dtr_pred))
print("Decision Tree RMSE:", np.sqrt(mean_squared_error(y_test,
dtr_pred)))

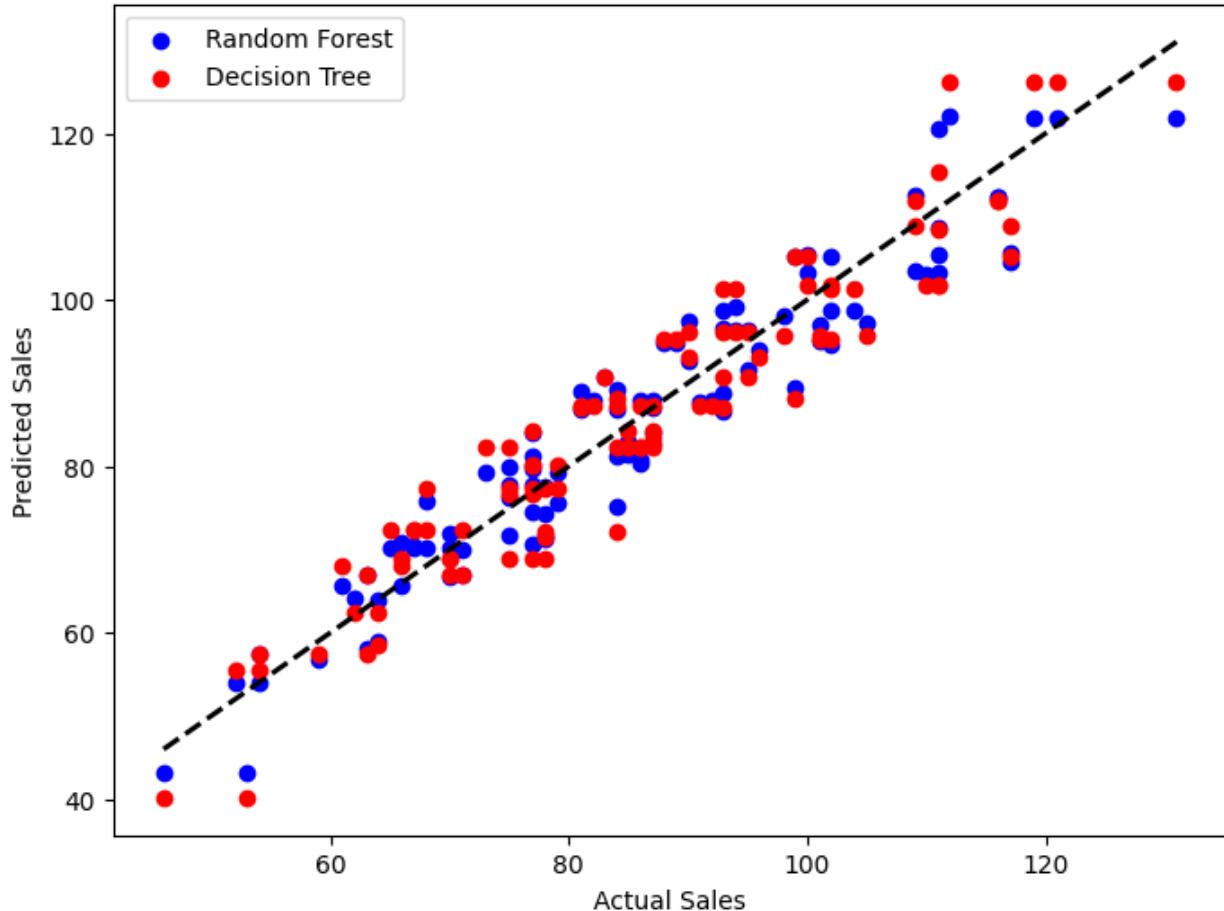
Random Forest R2: 0.9173629988719512
Random Forest MAE: 4.358646513624557
Random Forest RMSE: 5.110372167044106
Decision Tree R2: 0.9021360897168814
Decision Tree MAE: 4.6441429116268935
Decision Tree RMSE: 5.561302738480475

import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))
plt.scatter(y_test, rfr_pred, color='blue', label='Random Forest')
plt.scatter(y_test, dtr_pred, color='red', label='Decision Tree')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'k--', lw=2)
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs Predicted Jigarthanda Sales')
plt.legend()
plt.show()

```

Actual vs Predicted Jigarthanda Sales



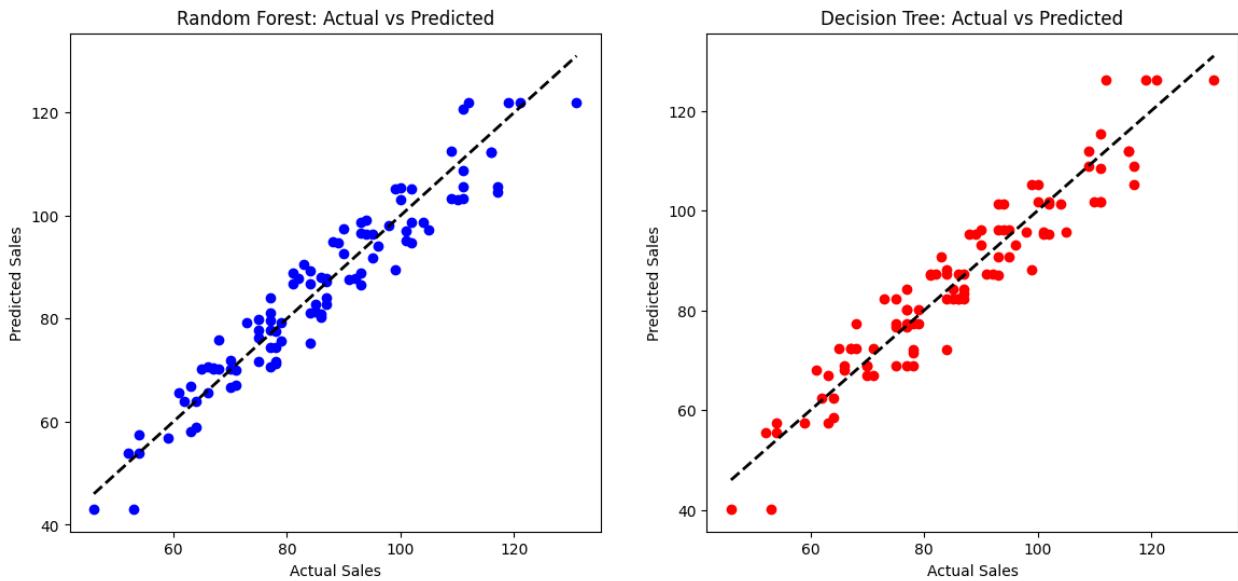
```
import matplotlib.pyplot as plt

fig, ax = plt.subplots(1, 2, figsize=(14,6))

# Random Forest
ax[0].scatter(y_test, rfr_pred, color='blue')
ax[0].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
ax[0].set_title('Random Forest: Actual vs Predicted')
ax[0].set_xlabel('Actual Sales')
ax[0].set_ylabel('Predicted Sales')

# Decision Tree
ax[1].scatter(y_test, dtr_pred, color='red')
ax[1].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
ax[1].set_title('Decision Tree: Actual vs Predicted')
ax[1].set_xlabel('Actual Sales')
ax[1].set_ylabel('Predicted Sales')
```

```
plt.show()
```



```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error
import numpy as np

models = ['Random Forest', 'Decision Tree']
predictions = [rfr_pred, dtr_pred]

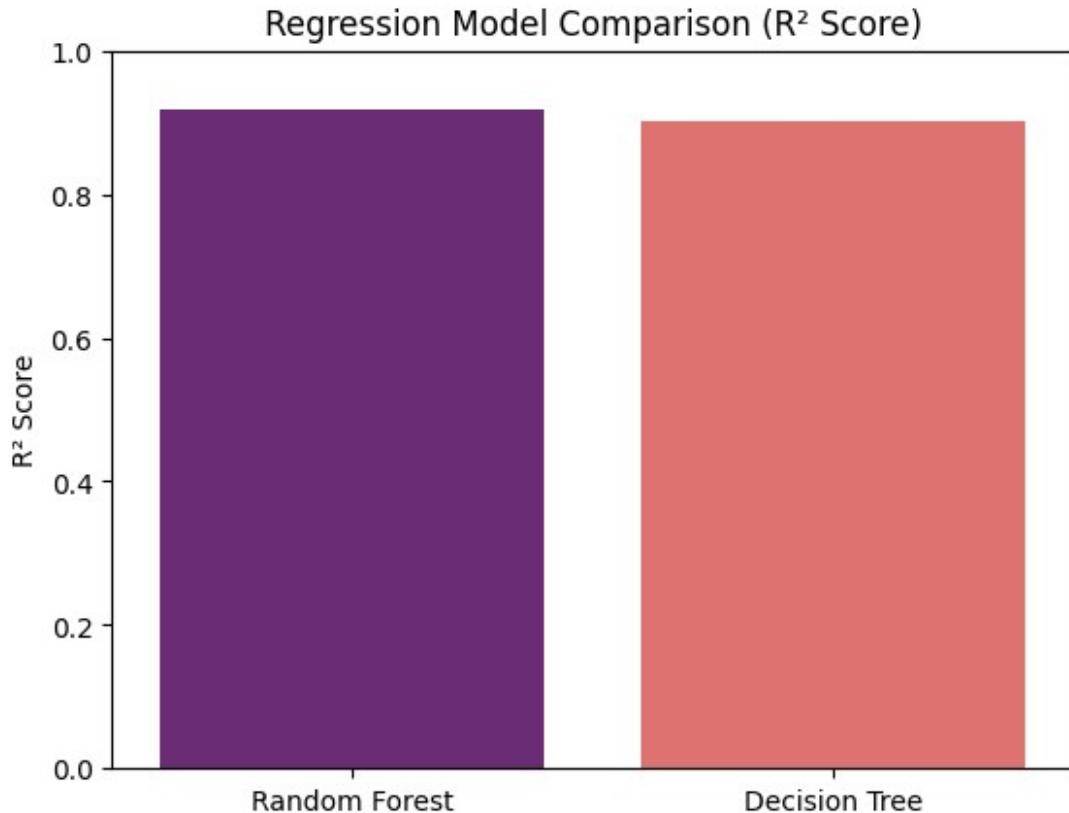
r2_scores = [r2_score(y_test, pred) for pred in predictions]

sns.barplot(x=models, y=r2_scores, palette='magma')
plt.title('Regression Model Comparison (R2 Score)')
plt.ylabel('R2 Score')
plt.ylim(0, 1)
plt.show()
```

```
C:\Users\PC\AppData\Local\Temp\ipykernel_6724\4062379592.py:11:
FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.
```

```
sns.barplot(x=models, y=r2_scores, palette='magma')
```



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error

le = LabelEncoder()
for col in ['Temperature', 'Weather', 'DayType']:
    df[col] = le.fit_transform(df[col])

X = df.drop('Jigarthanda_Sales', axis=1)
y = df['Jigarthanda_Sales']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=0)

rfr = RandomForestRegressor(
```

```

        criterion='squared_error',
        max_depth=8,
        min_samples_leaf=8,
        min_samples_split=2,
        random_state=42
    )

rfr.fit(X_train, y_train)
rfr_pred = rfr.predict(X_test)

param_grid_rfr = {
    'n_estimators': [100, 200, 300],
    'max_depth': [5, 10, 15, None],
    'min_samples_leaf': [1, 2, 4],
    'min_samples_split': [2, 5, 10],
    'criterion': ['squared_error', 'absolute_error'],
    'random_state': [0]
}

grid_rfr = GridSearchCV(estimator=RandomForestRegressor(),
param_grid=param_grid_rfr,
cv=5, n_jobs=-1, verbose=2, scoring='r2')

dtr = DecisionTreeRegressor(
    criterion='squared_error',
    max_depth=6,
    min_samples_leaf=6,
    min_samples_split=2,
    random_state=0
)

dtr.fit(X_train, y_train)
dtr_pred = dtr.predict(X_test)

param_grid_dtr = {
    'max_depth': [2, 4, 6, 8, 10],
    'min_samples_leaf': [1, 2, 4, 6],
    'min_samples_split': [2, 4, 6, 8],
    'criterion': ['squared_error', 'friedman_mse', 'absolute_error',
'poisson'],
    'random_state': [0, 42]
}

grid_dtr = GridSearchCV(estimator=DecisionTreeRegressor(),
param_grid=param_grid_dtr,
cv=5, n_jobs=-1, verbose=2, scoring='r2')

models = ['Random Forest', 'Decision Tree']

```

```

predictions = [rfr_pred, dtr_pred]

for name, pred in zip(models, predictions):
    print(f"\n{name} Metrics:")
    print("R2 Score:", r2_score(y_test, pred))
    print("MAE:", mean_absolute_error(y_test, pred))
    print("RMSE:", np.sqrt(mean_squared_error(y_test, pred)))

fig, ax = plt.subplots(1, 2, figsize=(14,6))

ax[0].scatter(y_test, rfr_pred, color='blue')
ax[0].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
           'k--', lw=2)
ax[0].set_title('Random Forest: Actual vs Predicted')
ax[0].set_xlabel('Actual Sales')
ax[0].set_ylabel('Predicted Sales')

ax[1].scatter(y_test, dtr_pred, color='red')
ax[1].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
           'k--', lw=2)
ax[1].set_title('Decision Tree: Actual vs Predicted')
ax[1].set_xlabel('Actual Sales')
ax[1].set_ylabel('Predicted Sales')

plt.show()

feat_importances = pd.Series(rfr.feature_importances_,
                             index=X_train.columns)
feat_importances.sort_values().plot(kind='barh', figsize=(8,6),
                                     color='teal', title='Random Forest Feature Importance')
plt.show()

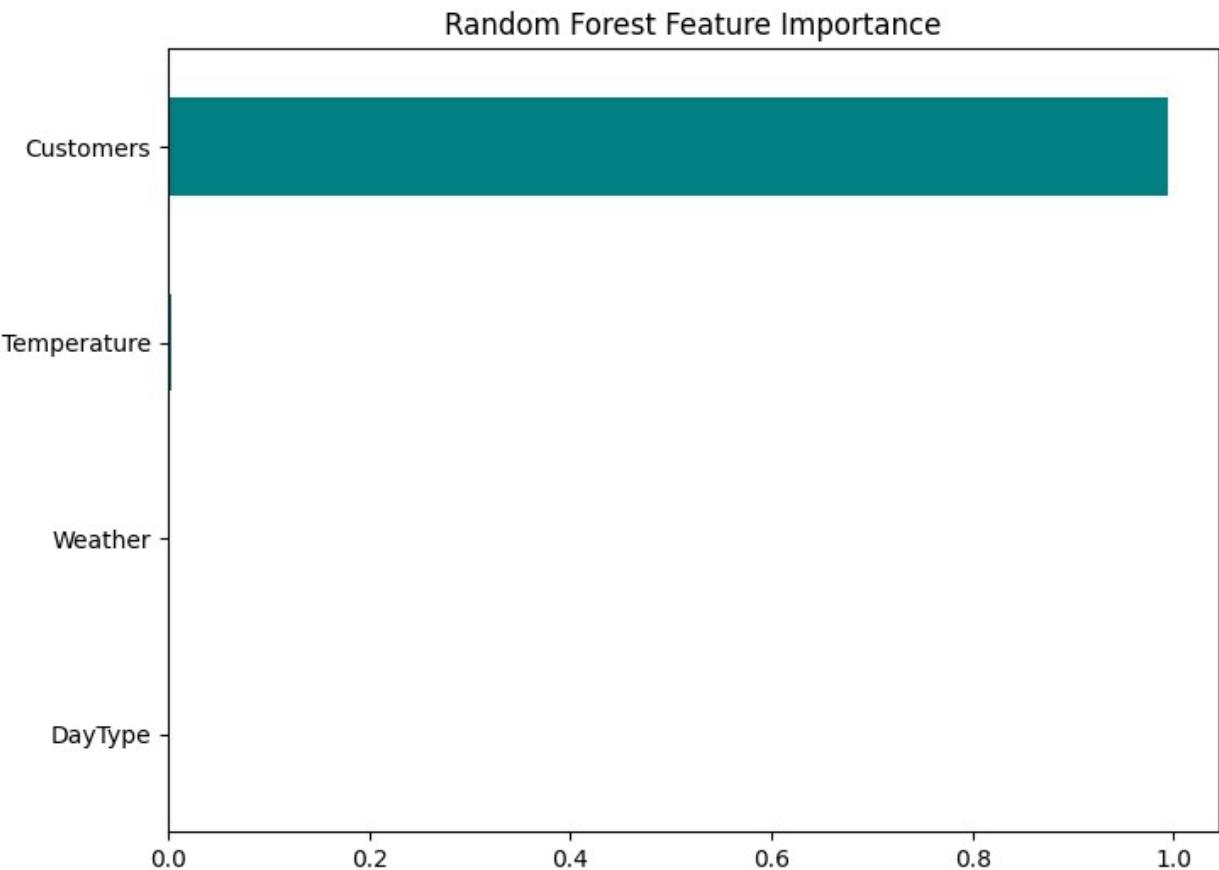
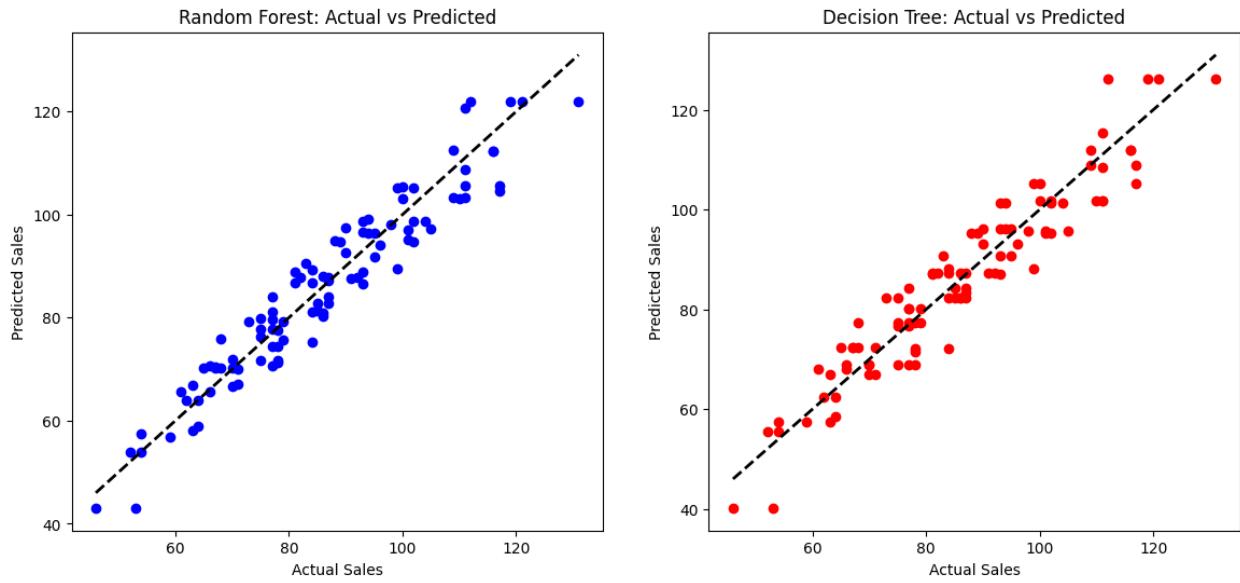
sns.histplot(y_test - rfr_pred, kde=True, color='blue', label='Random Forest Residuals')
sns.histplot(y_test - dtr_pred, kde=True, color='red', label='Decision Tree Residuals')
plt.title('Residual Distribution')
plt.xlabel('Residuals (Actual - Predicted)')
plt.legend()
plt.show()

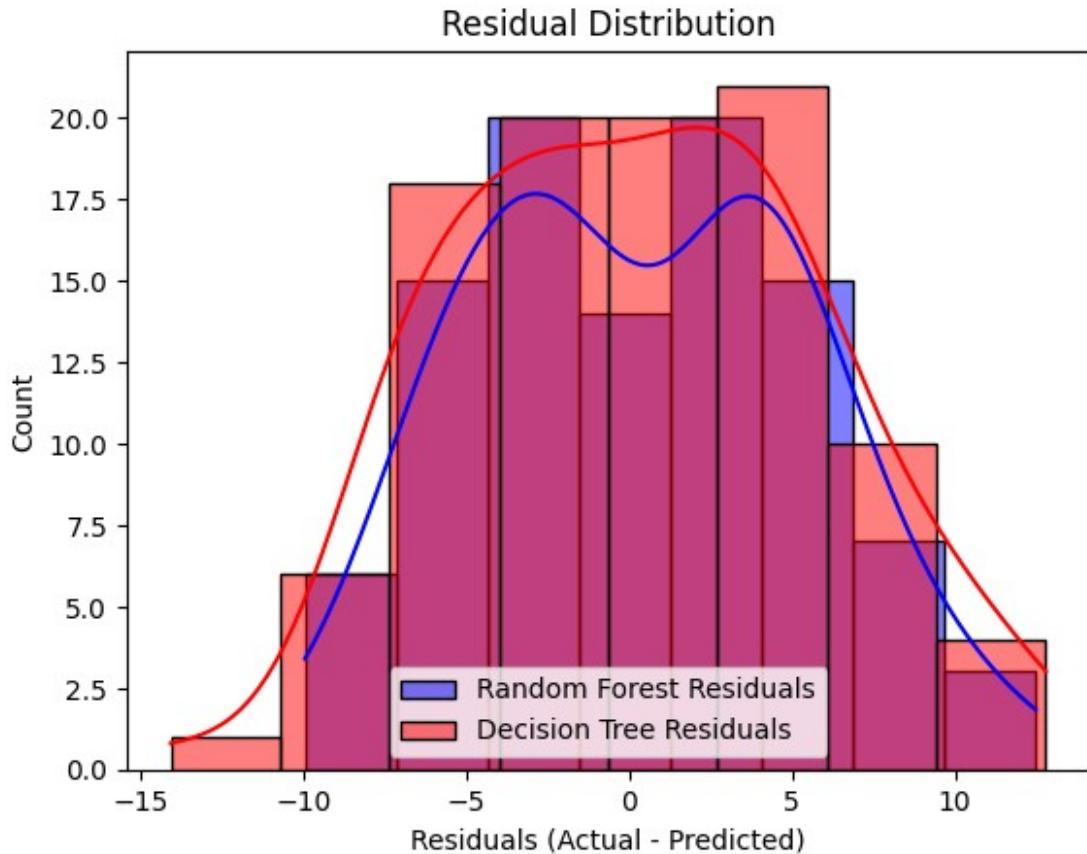
```

Random Forest Metrics:  
R2 Score: 0.9173629988719512  
MAE: 4.358646513624557  
RMSE: 5.110372167044106

Decision Tree Metrics:

R2 Score: 0.9021360897168814  
MAE: 4.6441429116268935  
RMSE: 5.561302738480475





## Summary

Your sales mainly increase on hot and sunny days, especially during weekends. Higher temperature and more customer footfall directly boost Jigarthanda sales. Afternoon hours (1 PM–4 PM) generally show stronger demand. Our Random Forest model predicts your daily sales with 91.7% accuracy, helping you plan staffing and inventory. To maximize sales, increase stock on hot days, weekends, and peak hours while running offers during slower periods.